

The 16S Workshop

Claire Duvallet

Thomas Gurry, Scott Olesen

March 9, 2016

after this workshop, I hope you will understand:

the steps involved in processing 16S data

what the AWS pipeline can do for you

how to troubleshoot some pipeline troubles

and 16S processing in general

where to find resources for help

mostly your eyes and your brain

the AWS pipeline is not trying to be qiime

AWS pipeline will one day be much more than qiime

different data types

automatic analyses

having an in-house pipeline offers other advantages

looking inside the black box

adding functionalities

troubleshooting with humans

qiime is great and you should totally use it if you want

but make sure you're using it wisely

“4” “easy” “steps” with 16S data

0. figure out what data you have

your eyes and your brain

1. pre-process data

AWS pipeline

2. build OTU table

AWS pipeline

3. analyze

your eyes and your brain

step 0. know your data

How were your regions sequenced?

barcode – fwd primer – region – (rev primer)



What kind(s) of file(s) do you have?

```
@FGG2ZB301CBTBW  
TACACGATCTACCATGCTGCCTCCCGTAGGAGTTGGGCCGTGTCTCAGTCCC AATGTGGCCGGTCACCTCTCAGGTCGGCTATGGATCGTCGG  
+FGG2ZB301CBTBW  
@GFFFFFFFFFFFFFIIIGGGIIIIIIIIFFFFFFFF  
@FGG2ZB301D7B09  
GTGGCGATACCATGCTGCCTCCCGTAGGAGTCTGGTCCGTGTCTCAGTACCAGTG TGGGGTTAACCTCTCAGTCCCCCTATGTATCGTCGC  
+FGG2ZB301D7B09  
FFFFFFFFFFFFFFFFIIIIIIIIFFFFFFFFCC5555==FFFFFFFFFG????GFFFFFFFF
```

```
>SWDL50_1
CTGGGCCGTGTCTCAGTCCCAATGTGGCCGTTCAACCTCTCAGGCCGGCTATGGATCGTCGCCTTGGTAGGCCGTTACCCACCAAG
>SWDL27_1
TTGGGCCGTGTCTCAGTCCCAATGTGGCCGTTCAACCTCTCAGGCCGGCTACTGATCGTCGCCTTGGTGGGCCGTTACCCGCCAAG
>SWDL50_2
CTGGGCCGTGTCTCAGTCCCAATGTGGCCGTTCAACCTCTCAGTCCGGCTACCGATCGTCGCCTTGGTGGGCCATTACCTCACCAAG
>SWDL15_2
CTGGACCGTGTCTCAGTTCAGTGTGGCCGTTCACTCTCAGACCGGCTACTGATCGTAGGTTTGGTGGGCCGTTACCTCACCAAG
>SWDL36_2
TTGGGCCGTGTCTCAGTCCCAATGTGGCCGGTCAACCTCTCAGGCCGGCTACTGATCGTCGCCTTGGTGGGCCGTTACCTCACCAAG
```

step 0. know your data

How were your regions sequenced?

barcode – fwd primer – region – (rev primer)



What kind(s) of file(s) do you have?

fastq

[illegible]

fasta

```
>SWDL50_1
CTGGGCCGTGTCTCAGTCCCAATGTGGCCGTTCAACCTCTCAGGCCGGCTATGGATCGTCGGCTTGGTAGGCCGTTACCCACCAAG
>SWDL27_1
TTGGGCCGTGTCTCAGTCCCAATGTGGCCGTTCAACCTCTCAGGCCGGCTACTGATCGTCGCCTTGGTGGGCCGTTACCCGCCAAG
>SWDL50_2
CTGGGCCGTGTCTCAGTCCCAATGTGGCCGTTCAACCTCTCAGTCCGGCTACCGATCGTCGCCTTGGTGGGCCATTACCTACCAAG
>SWDL15_2
CTGGACCGTGTCTCAGTTCAGTGTGGCCGTTATCTCTCAGACCGGCTACTGATCGTAGGTTTGGTGGGCCGTTACCTACCAAG
>SWDL36_2
TTGGGCCGTGTCTCAGTCCCAATGTGGCCGGTCAACCTCTCAGGCCGGCTACTGATCGTCGCCTTGGTGGGCCGTTACCTACCAAG
```

step 0. understand your data

single end reads? paired end? already merged?

if not merged, you need to do this step yourself

one single fastq? one fastq file per sample?

if downloading data: already a fasta?

what about the sequences?

barcodes removed?

primers removed?

length trimmed already?

every sequencing center does it differently!!

use your eyes and your brain

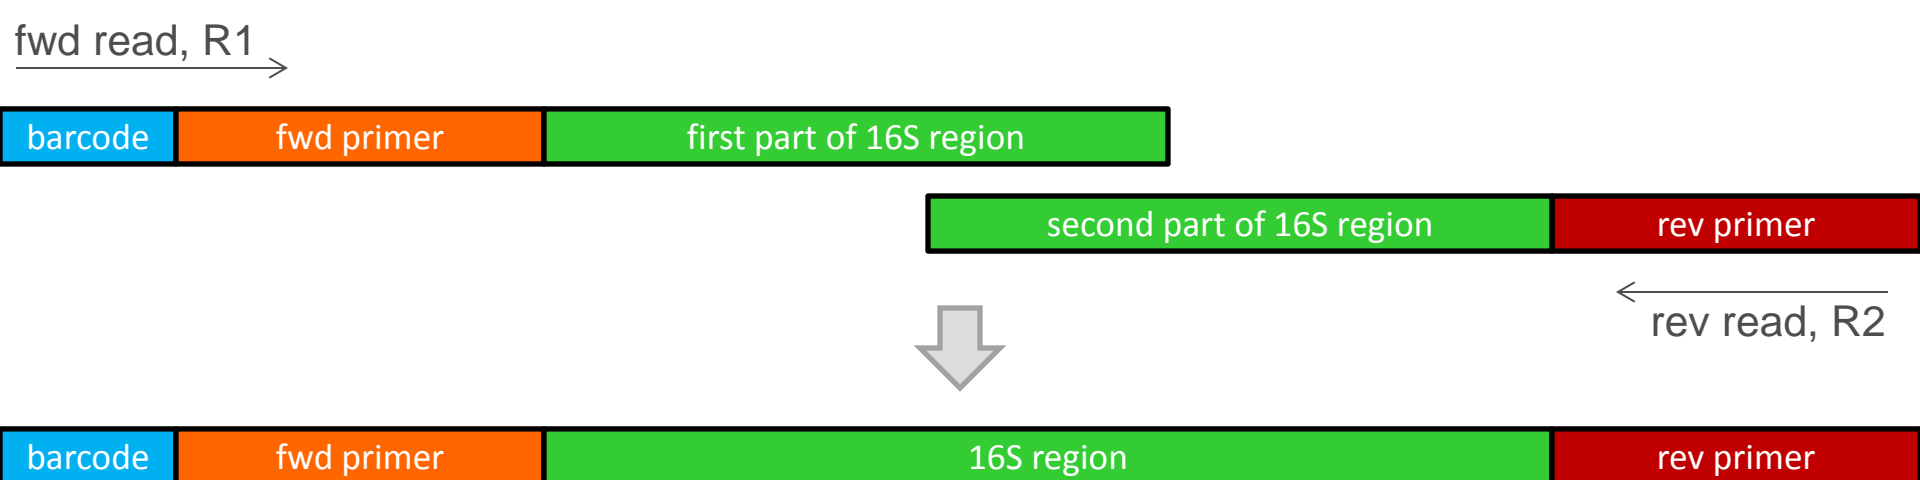
step 0. paired-end sequencing will require merging

merge forward and reverse reads

can use usearch or Scott's script

note: barcodes may be in a separate index file

AWS pipeline does not do this step



“4” “easy” “steps” for 16S data

0. figure out what data you have

your eyes and your brain

1. pre-process data

AWS pipeline

2. build OTU table

AWS pipeline

3. analyze

your eyes and your brain

step 1. pre-processing



Raw data options:

- one fastq or multiple fastq files
- with or without barcodes in sequences
- with or without primers in sequences
- all different lengths, possibly low quality
- can also take in a fasta

Outputs:

- dereplicated fasta: all of the unique sequences
- dereplication map: counts of each sequence per sample

step 1. pre-processing with AWS pipeline


which steps are done depends on inputs

talk to pipeline via `summary_file.txt`

AWS pipeline calls a medley of scripts

lab scripts (`2.split_by_barcodes.py`, etc)

usearch functions



step 1.1: need to split by barcodes if fastqs aren't already demultiplexed

if you only have one fastq file with all samples in it, need to split by barcodes

basically put sample IDs in the fastq sequence IDs

sequencing centers often demultiplex fastqs for you

i.e. you have one file per sample



step 1.1: split by barcodes

barcodes can be in multiple places


beginning of sequence, sequence ID, index file
this can be specified in `summary_file.txt`

pipeline calls `2.split_by_barcodes.py`

or simply renames sequences with sample IDs
if fastq is already de-multiplexed



step 1.1: demultiplexing also trims off
barcode and relabels sequences



xaa

or original fastq file name

xaa.sb

sample IDs in seq IDs
barcodes removed
no seqs with barcode mismatches

```
@FGG2ZB301D4YMD  
TAGCCTCTCTGCCATGCTGCCTCCCGTAGGAGTCTGGGCCGTGTCTCAGTCCC  
+FGG2ZB301D4YMD  
FFFFFFFFFFFFFFFFIIIIIIIIIIIIIIIIIIIIFFFFFFFFFFF  
@FGG2ZB301DRL30  
TAACCTCTGATGCCATGCTGCCTCCCGTAGGAGTCTGGGCCGTGTCTCAGTCCC  
+FGG2ZB301DRL30  
FFFFFFFFFFFFFFFFIIIIIIIIIIIIIIIIIIIIFFFFFFFFFFF  
@FGG2ZB301D5A04  
TAGCACACCTATCATGCTGCCTCCCGTAGGAGTCTGGGCCGTGTCTCAGTCCC  
+FGG2ZB301D5A04  
FFFFFFFFFFFFFFFFIIIIIIIIIIIIIIIIIIIIFFFFFFFFFFF
```



```
@SWDL48_1  
CATGCTGCCTCCCGTAGGAGTCTGGGCCGTGTCTCAGTCCAATGTGGCCGTTCT  
+FGG2ZB301D4YMD  
FFFFFFFFIIIIIIIIIIIIIIIIIIIIIIIIIFFFFFFFFFFFFFFFFFFFFFFFF  
@SWDL36_1  
CATGCTGCCTCCCGTAGGAGTCTGGGCCGTGTCTCAGTCCAATGTGGCCGTTCT  
+FGG2ZB301DRL30  
FFFFFFFFIIIIIIIIIIIIIIIIIIIIIIIIIFFFFFFFFFFFFFFFFFFFFFFFF  
@SWDL46_1  
CATGCTGCCTCCCGTAGGAGTCTGGGCCGTGTCTCAGTCCAATGTGGCCGTTCT  
+FGG2ZB301D5A04  
FFFFFFFFIIIIIIIIIIIIIIIIIIIIIIIIIFFFFFFFFFFFFFFFFFFFFFFFF
```

step 1.2: remove primers if they're still in sequences

should manually check if primers are in seqs

pipeline calls wrapper to `1.remove_primers.py`

checks for primers in forward and reverse complement sequence



step 1.3: quality trimming

you need to figure out quality score encoding

use `usearch -fastq_chars` to make a guess

specify in `ASCII_ENCODING` in summary file

AWS quality trims with `usearch -fastq_filter`

truncates at first base with $p(\text{error}) > 0.003$

default `Q=25`, can specify in summary file

what kind of quality filtering do you prefer?



step 1.4: length trimming

sequences must be same length for later steps

i.e. alignment and clustering

AWS length trims with `usearch -fastx_truncate`

AWS default is 101, specify in `TRIM_LENGTH`

results are in `*.raw_trimmed.fasta`

has all of the reads for all samples

all same length, no barcodes or primers

sample IDs are in sequence headers

we're almost there!



step 1.5: dereplication

keeps only the unique sequences

in AWS these are in *.raw_dereplicated.fasta

and does “provenancing”

i.e. maps how many times each sequence
appeared in each sample

in AWS this file is *.dereplication_map

pipeline calls `3.dereplication.py`



step 1.5: dereplication documentation is confusing but outputs are recognizable

dereplication map

```
1 NASH31:13
2 NASH34:4 NASH67:1 NASH11:1 NASH44:1 NASH57:3 NASH55:1
3 NASH26:43 NASH44:1 NASH37:1 NASH39:1 NASH56:2
4 NASH17:34 NASH73:21 NASH72:11 NASH74:14 NASH11:19 NASH10:20
5 NASH26:2 NASH27:6 NASH22:36 NASH32:1 NASH17:1 NASH72:2 NASH44:1
6 NASH58:1 NASH13:1 NASH55:8 NASH21:2
7 NASH37:37
8 NASH62:3 NASH7:5 NASH36:1 NASH3:1 NASH71:10 NASH63:1 NASH61:5
9 NASH9:1 NASH23:1 NASH30:1 NASH52:3 NASH51:1 NASH40:1 NASH19:1
10 NASH16:10 NASH72:17 NASH74:10 NASH52:1 NASH51:1 NASH50:1 NASH17:1
11 NASH17:4 NASH16:12 NASH72:33 NASH52:46 NASH51:7 NASH56:2 NASH17:1
12 NASH62:1 NASH22:1 NASH37:1 NASH4:1 NASH3:1 NASH17:2 NASH72:3
13 NASH17:1 NASH6:1 NASH37:1 NASH4:18 NASH3:1 NASH1:28 NASH21:2
14 NASH62:1 NASH38:5 NASH61:1 NASH.CR:1 NASH30:1 NASH2:1 NASH72:1
```

sequence 1 occurred 13 times in sample NASH31

sequence 2 occurred 4 times in sample NASH34, 1 time in NASH67, 1 time in NASH11, ...

sequence 3 occurred 43 times in NASH26, 1 time in NASH37, 1 time in NASH39, and 2 times in NASH6

dereplicated fasta

```
>988;size=86698
TACGGAGGATCCGAGCGTTATCCGGATTTATTGGGTTTAAAGGGAGCGTAGATGGATGTTTAAGTC
>1993;size=68808
TGTTTGCTCCCCACGCTTTCGAGCCTCAACGTCAGTTACCGTCCAGTAAGCCGCCTTCGCCACTGG
>1487;size=67810
TGTTTGATACCCACACTTTCGAGCCTCAATGTCAGTTGCAGCTTAGCAGGCTGCCTTCGCAATCGG
>147;size=37154
TACGGAAGGTCCGGGCGTTATCCGGATTTATTGGGTTTAAAGGGAGTGATAGGCGGCCTGTTAAGCG
>383;size=35730
TGTTTGCTACCCACACTTTCGAGCCTCAGCGTCAGTTGGTGCCAGTAGGCCGCCTTCGCCACTGG
>802;size=34832
TACGGAAGGTCCGGGCGTTATCCGGATTTATTGGGTTTAAAGGGAGCGTAGGCCGAGATTAAGCG
>2367;size=34134
TGTTTGCTCCCCACGCTTTCGTGCCTCAGCGTCAGTTCAAGTCCAGAAAGTCGCCTTCGCCACCGG
>2215;size=26354
TACGTAAGGGGCGAGCGTTGTCCGGAATTATTGGGCGTAAAGAGTGCGTAGGCCGCAAAATTAAGTC
```

unique sequences only
sequences sorted by counts

step 1: dereplicated fasta and map are the end goal of pre-processing

pipeline handles many possible inputs

one fastq or multiple fastq files

with or without barcodes in sequences

with or without primers in sequences

all different lengths, possibly low quality

can also take in a fasta


all of this gets specified in `summary_file.txt`

if there are no primers/barcodes, put None

for other options, read documentation



step 1: AWS pipeline saves intermediate files, identified with suffixes



```
xaa
xaa.sb
xaa.sb.pt
xaa.sb.pt.qt
xaa.sb.pt.qt.lt
xaa.sb.pt.qt.lt.fasta
xab
xab.sb
xab.sb.pt
xab.sb.pt.qt
xab.sb.pt.qt.lt
xab.sb.pt.qt.lt.fasta
```

so you can use your eyes and your brain

processing: eyes and brain practice

datasets on Alm lab node

what data do we have?

fastq or fasta?

de-multiplexed already?

primers in there?

need to merge?

trimmed yet?

processing: eyes and brain practice

datasets on workshop node

```
ssh -i 16S_workshop.pem  
ubuntu@52.70..179.175
```

for datasets 7-10: what happened?

where would you start de-bugging?

where can you find more info about what went wrong?

where did it fail? what went wrong?

dataset7

dataset8

```
Traceback (most recent call last):
  File "/home/ubuntu/scripts/2.split_by_barcode.py", line 177, in <module>
    run()
  File "/home/ubuntu/scripts/2.split_by_barcode.py", line 124, in run
    b2s = parse_barcode_file(args.b, format=args.B, rc=args.rc) # barcodes to samples
  File "/home/ubuntu/scripts/2.split_by_barcode.py", line 51, in parse_barcode_file
    [s,b] = line.rstrip().split()
ValueError: too many values to unpack
```

dataset9

expected 128 samples, got 56
have intermediate files

dataset10

USEARCH broke
have intermediate files

what went wrong? (hints)

dataset7

no useful errors sry

```
00:00 99Mb 0.1% Filtering  
WARNING: Option -fastq_truncqual ignored  
  
WARNING: Option -fastq_truncqual ignored  
ubuntu@ip-10-0-1-131:~/datasets/ra_littman_2013$
```

dataset8

```
Traceback (most recent call last):  
  File "/home/ubuntu/scripts/2.split_by_barcodes.py", line 177, in <module>  
    run()  
  File "/home/ubuntu/scripts/2.split_by_barcodes.py", line 124, in run  
    b2s = parse_barcodes_file(args.b, format=args.B, rc=args.rc) # barcodes to samples  
  File "/home/ubuntu/scripts/2.split_by_barcodes.py", line 51, in parse_barcodes_file  
    [s,b] = line.rstrip().split()  
ValueError: too many values to unpack
```

dataset9

```
fastq_to_fasta: Premature End-Of-File (filename = 'xak.qt.lt')  
fastq_to_fasta: Premature End-Of-File (filename = 'xaf.qt.lt')  
fastq_to_fasta: Premature End-Of-File (filename = 'xal.qt.lt')  
fastq_to_fasta: Premature End-Of-File (filename = 'xam.qt.lt')  
fastq_to_fasta: Premature End-Of-File (filename = 'xan.qt.lt')  
fastq_to_fasta: Premature End-Of-File (filename = 'xao.qt.lt')  
ubuntu@ip-10-0-1-131:~/datasets/ibd_engstrand_2009$ |
```

dataset10

“4” “easy” “steps” for 16S data

0. figure out what data you have

your eyes and your brain

1. pre-process data

AWS pipeline

2. build OTU table

AWS pipeline

3. analyze

your eyes and your brain

step 1. pre-processing



step 2. build OTU table



step 2: building an OTU table is essentially combining dictionaries

an OTU table is a nested dictionary

it shows the counts for each OTU in each sample

```
{sample1: {OTU1: counts, OTU2: counts}...}
```

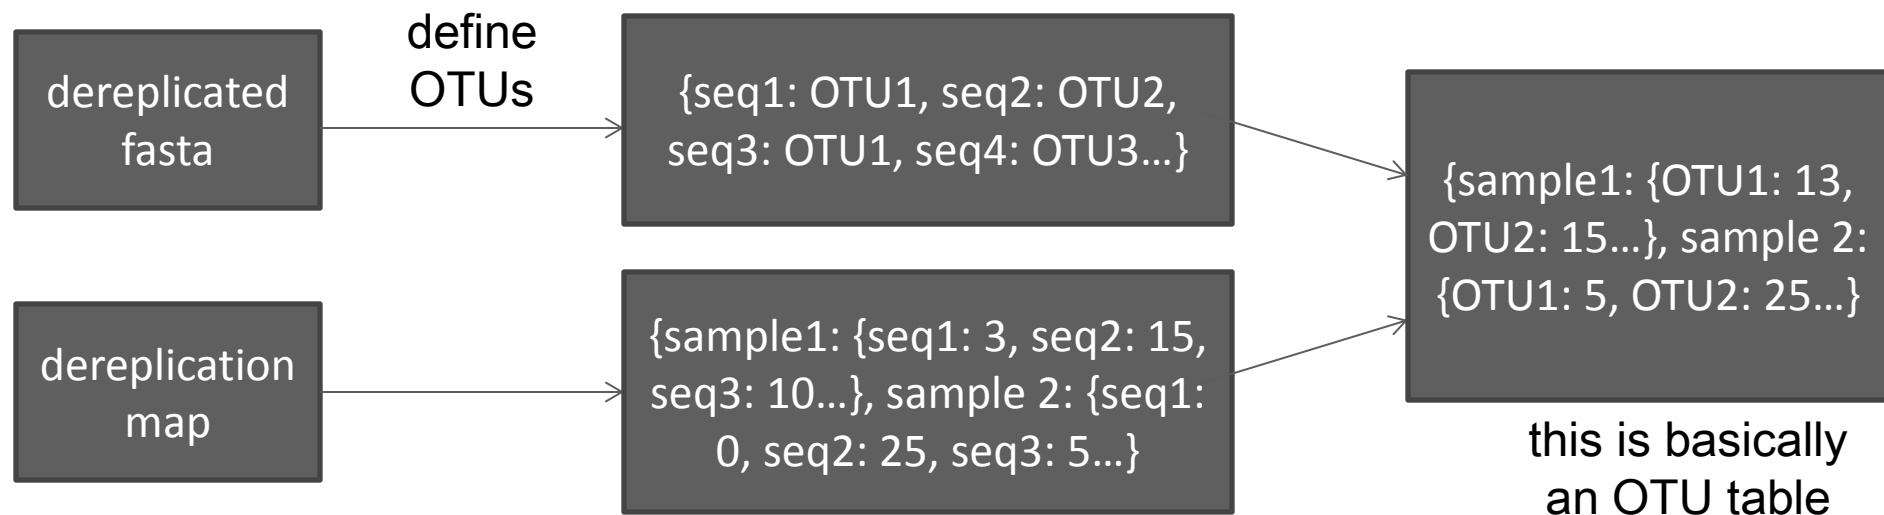
so we first need to map:

sequence ID → OTU ID for each original seq
sequence ID → counts for each sample

step 2: building an OTU table is essentially combining dictionaries

dereplicated fasta is used to map sequences to OTUs

dereplication map already contains the map of sequence counts to samples



step 2: there are 3 categories of methods for defining OTUs

de novo

make OTUs based only on sequences (ATCG)
no database

closed reference

map sequences to a database to define OTUs
throw out anything that doesn't match
most commonly: map to Green Genes

open reference

map sequences to a database
then do de novo clustering on leftover sequences
and add to the mapped sequences

step 2: denovo OTU calling approaches

similarity based OTU clustering

simplest and most common

usearch -cluster_otus

also removes chimeras

oligotyping

uses information theory to separate meaningful sequences

distribution-based clustering (DBC)

uses both similarity and distribution of sequences to call OTUs

DADA2

seems awesome, only works for Illumina data

if not using usearch, you need to remove chimeras!

at some point in the processing pipeline

step 2: all OTU calling approaches have pros and cons

denovo

pro: no need for database

con: need to do more work for taxonomic info

con: can't compare across datasets

closed reference

pro: get taxonomic info for cheap

pro: might be able to compare across datasets

con: database-dependent

con: how clustering method deals with ties is important

con: throws out all unmatched sequences

open reference

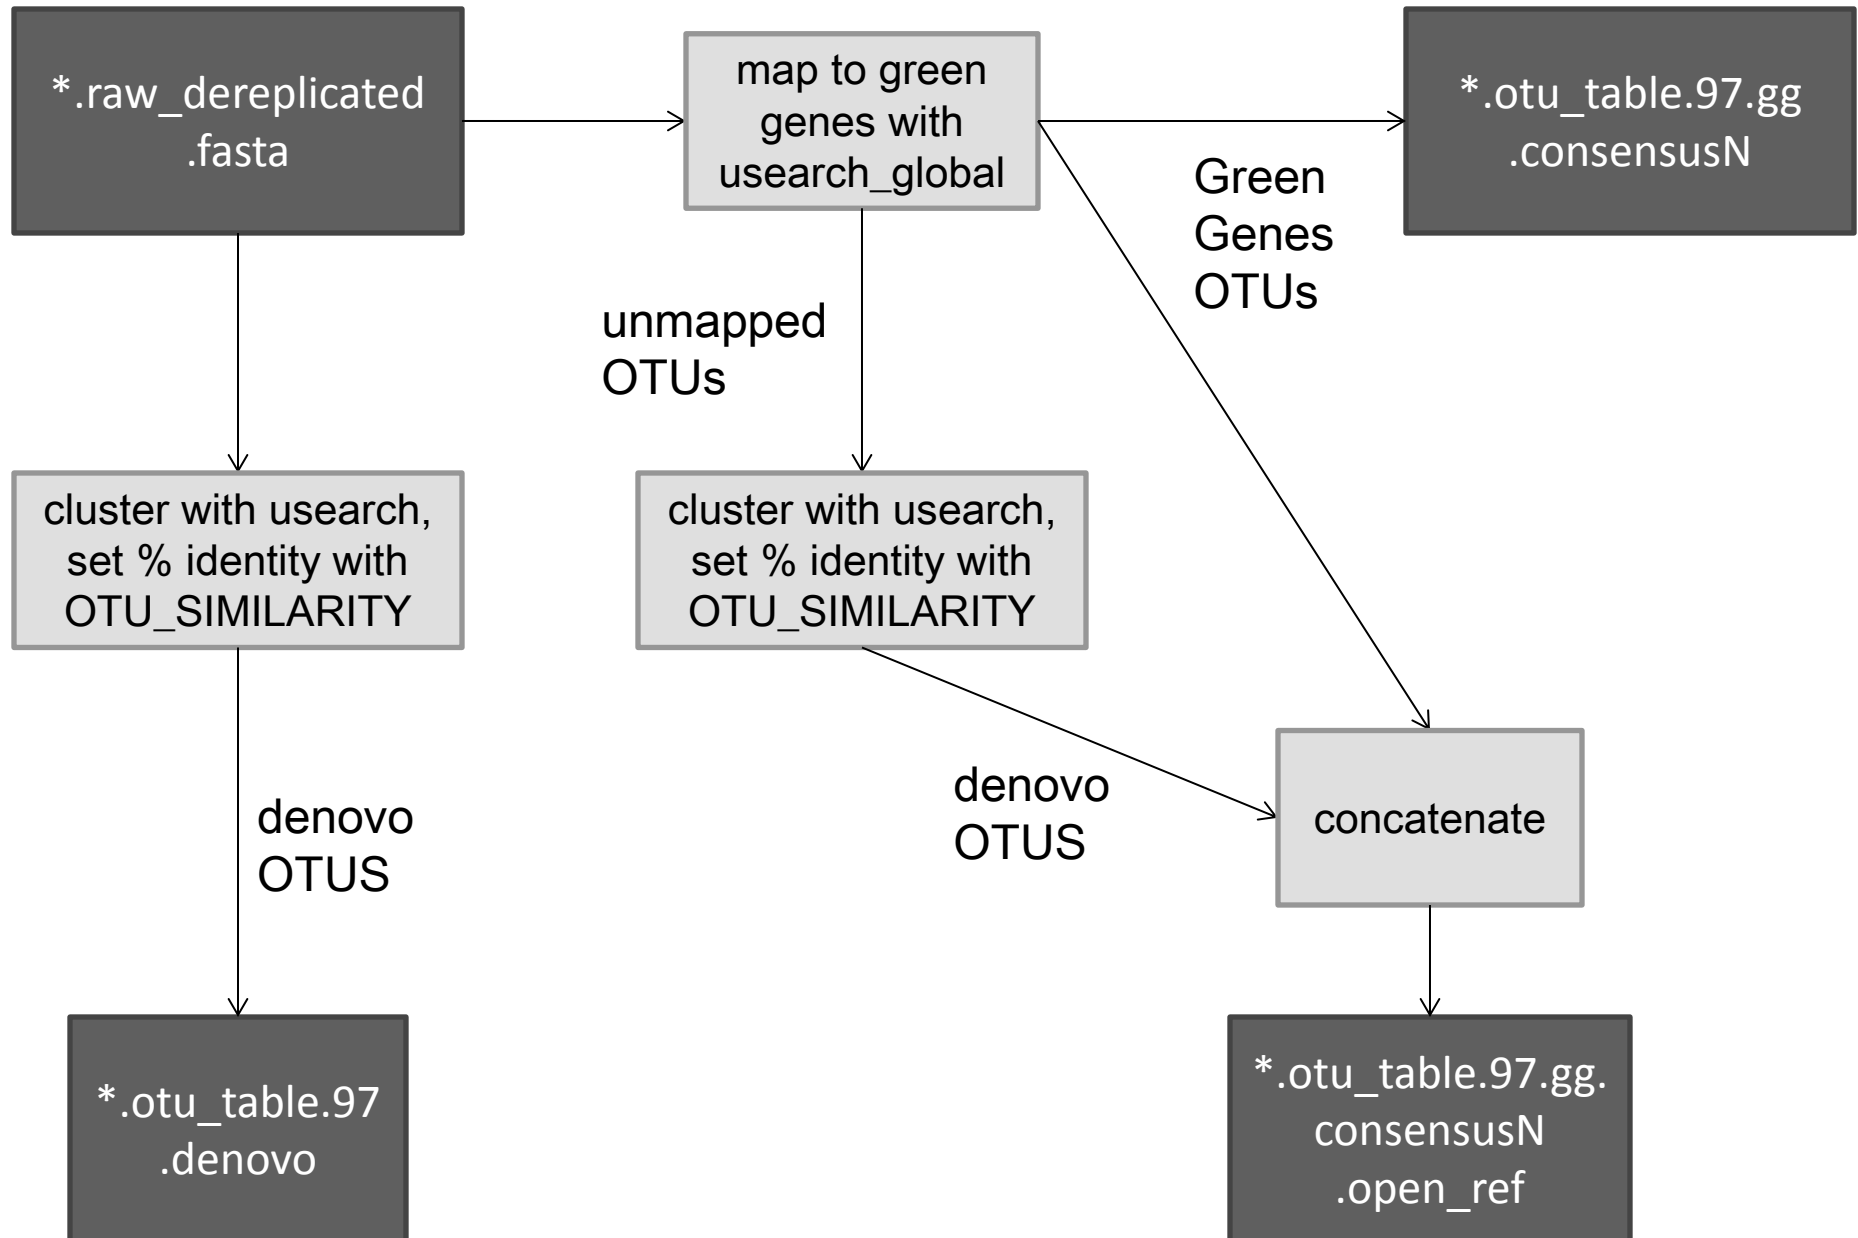
pro: you don't throw out your unknown sequences

pro: have taxonomic info for some OTUs

cons: same disadvantages as denovo and closed

cons: is the distance metric really the same?

step 2: AWS pipeline spits out heaps of files



step 2: but wait there's more!

*.otu_table.97
.denovo

find unique seqs
within each
denovo OTU

assign Latin names,
set confidence with
RDP_CUTOFF

*.otu_table.97
.denovo_oligotypes

*.otu_table.97
.denovo
.rdp_assigned

*.otu_seqs.97.fasta

contains representative
sequences for the denovo
OTUs

*.otu_table.97.gg
.consensusN

GreenGenes assignment
obtained by taking the
consensus of the top N hits
when aligned to the
GreenGenes database

step 3: analyze data – sanity checks

sanity checks


what's in your OTU table

number of reads per sample


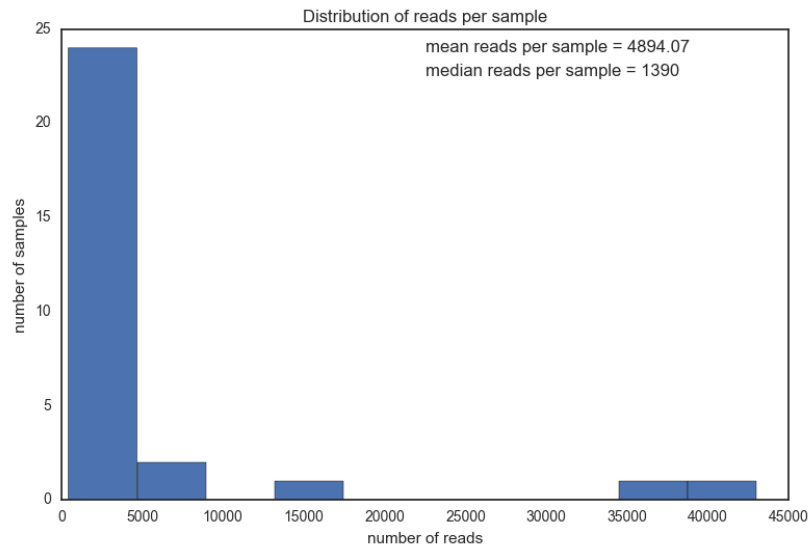
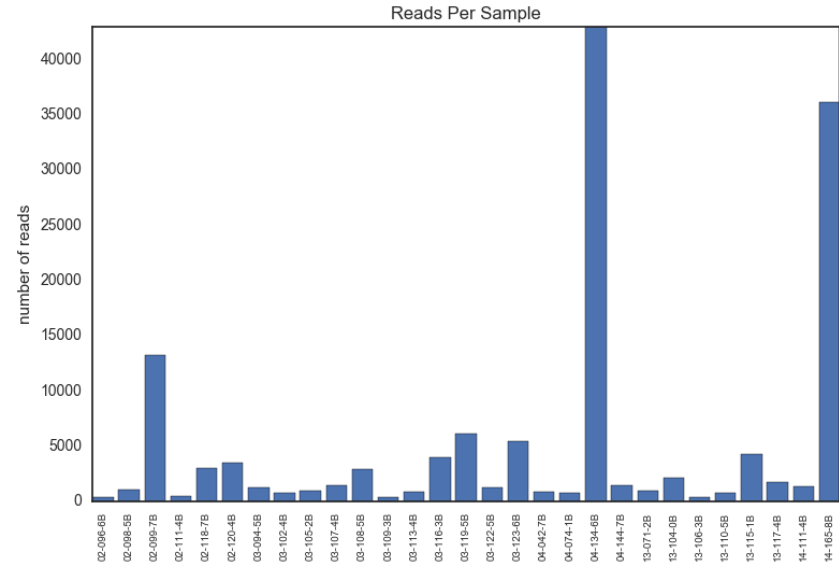
distribution of OTUs across samples

heatmap of abundances

```
In [32]: otu_table.shape  
Out[32]: (583, 1)
```



step 3: low reads per sample can be bad but can also be okay




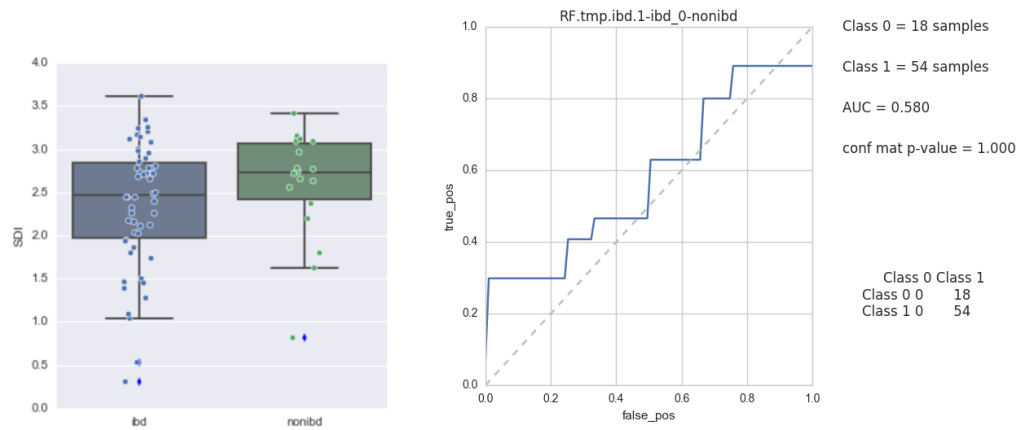

step 3: some places to start with analysis

alpha diversity

PCA

univariate comparisons

basic classifiers



“4” “easy” “steps” for 16S data

0. figure out what data you have

your eyes and your brain

1. pre-process data

AWS pipeline



2. build OTU table

AWS pipeline



3. analyze

your eyes and your brain

using the pipeline is easy and looking
inside the black box is also easy

code currently on AWS and private github

Master.py calls raw2otu.py

reads in `summary_file.txt` for file
locations, parameter settings, etc

raw2otu.py does most of its work in the
preprocessing_16S.py module

many functions are wrappers to either usearch
functions or Alm lab scripts

there's a dropbox folder from 20.106
with problem sets and sample data

play around with the files

look at what input and output files look like
figure out the syntax

don't be limited by the code given

some steps are missing in instructions (i.e.
length and quality trim)

you can also play around with parameters

<https://www.dropbox.com/sh/enyftm9ut2r74ry/AAAZejnQlhUydVNsrCnWgrr-a?dl=0>

thanks!

